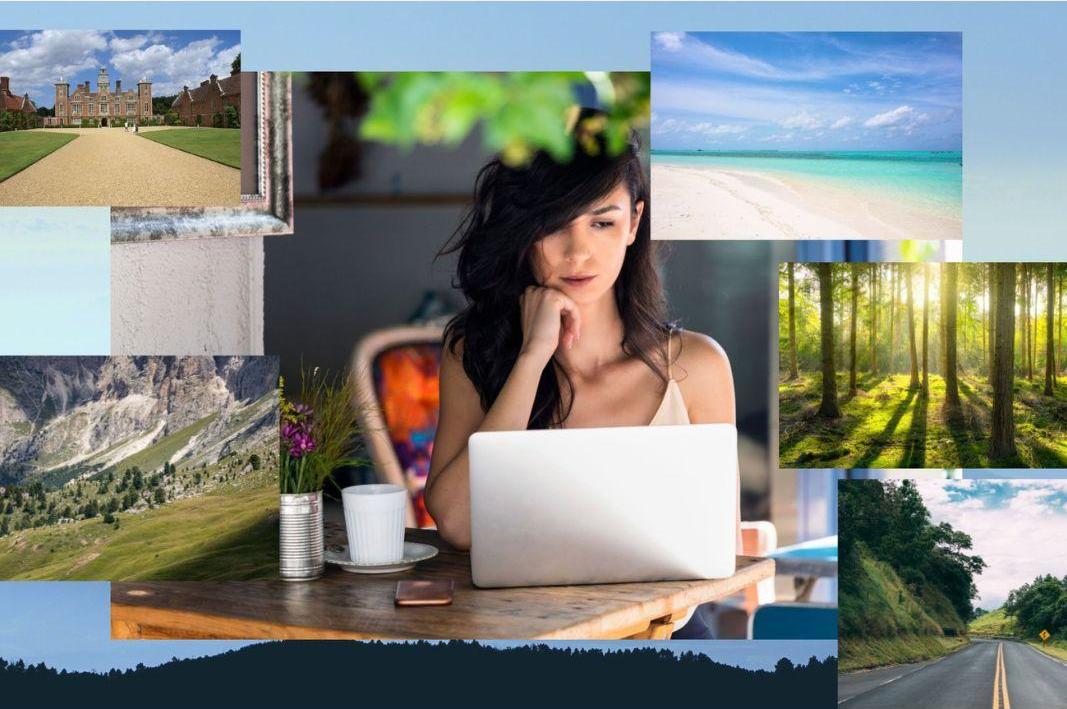




# SCENE CLASSIFICATION

DATASCI 207 FINAL PROJECT OF TEAM : KT NORTON, YURI KINAKIN & JUDD GALLARES  
PROF. NEDELINA TENEVA, PHD

DECEMBER 6, 2022



# Problem Motivation

We work for a photography travel blog website that allows its users to share, view and interact with other users' content and photography.

# Problem Motivation

We want to recommend blog articles to our users based on photos that they have liked, interacted with, or posted



→ C  kaggle.com/datasets/nitishabharathi/scene-classification

Search

Sign In Register

NITISHA - UPDATED 4 YEARS AGO

58 New Notebook Download (378 MB)

## Scene Classification

Contains ~25K images from a wide range of natural scenes



# Problem Statement

Data source: Kaggle.com

We will explore various techniques to classify images based on their natural scene or setting.

# Dataset Description

The data contains ~25k images from around the world of natural outdoor scenes, classified in one of 6 classes.



1. FORESTS

2. MOUNTAINS

3. STREETS

4. SEA

5. BUILDINGS

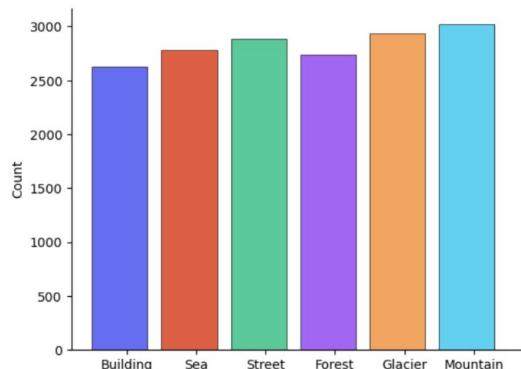
6. GLACIERS

# Dataset Description

<b>train</b>	A folder containing 24,335 of mostly 150x150px colored (rgb) jpegs	 - 1.jpg - 12.jpg - 123.jpg ...
<b>train.csv</b>	A csv file with 17,034 rows of labels each describing an image in <b>train</b>	 - Buildings - Forests - Mountains - Glaciers - Street - Sea
<b>test_WyRytb0.csv</b>	A csv file containing 7,301 rows of image names (strings) corresponding to a jpeg in <b>train</b> . The <b>train</b> and <b>test</b> csvs should have no duplicates.	 - "3.jpg" - "10.jpg" - "168.jpg" ...

# Exploratory Data Analysis

## Scene Classification EDA and Embeddings



### Dataset Summary

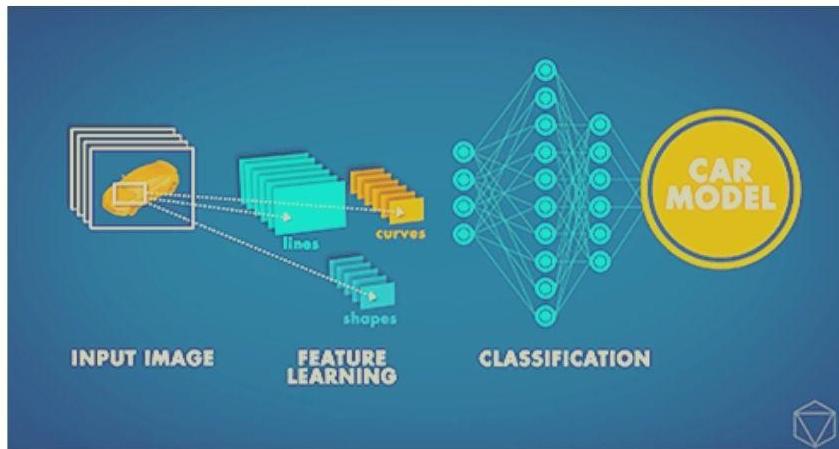
A total of 16,979 labelled images (150px x 150px x 3rgb channels) are available for training.

They are split into 6 different categories of approximately equal size.

Our baseline accuracy for each class is therefore ~17%.

# Solution/Approach Description

## Convolutional Neural Network (CNN) 2D - RGB and 1D - Grayscale

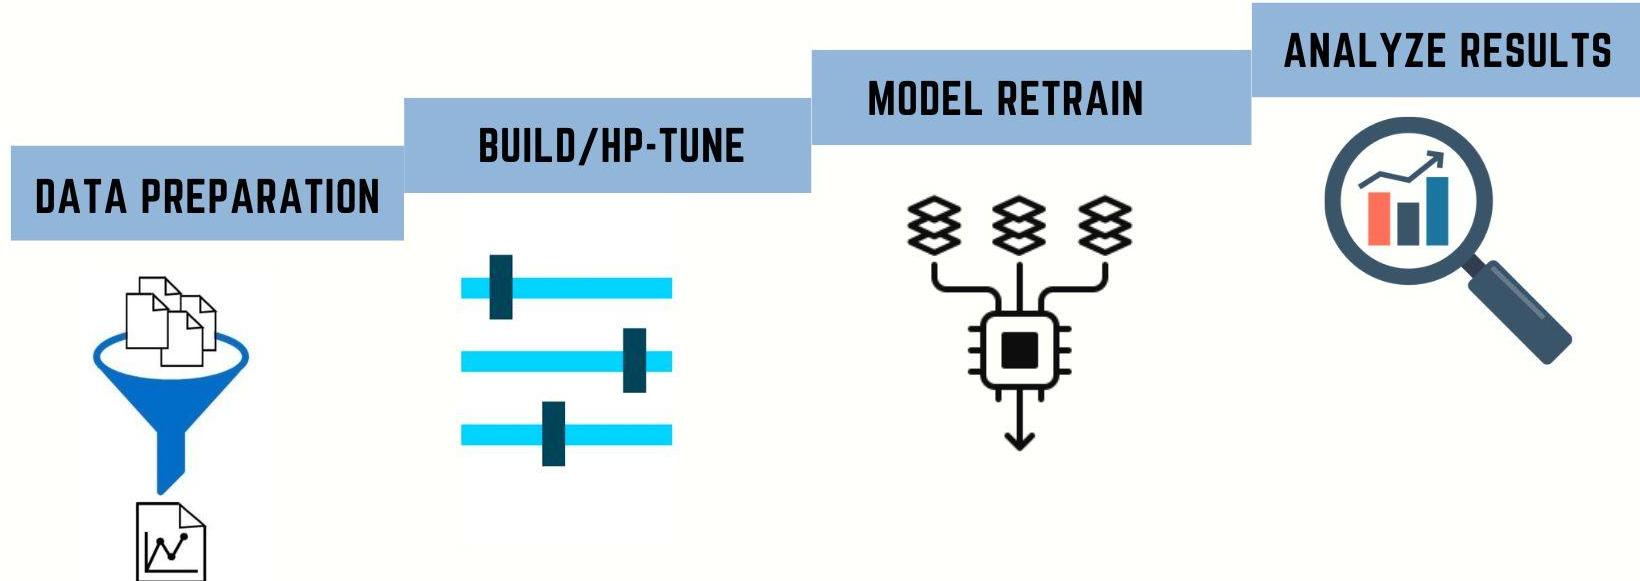


Our initial analysis utilizing classical machine learning techniques result in a relatively low and imbalance classification accuracy. Our team then tested the approaches below to see which one produced the best results.

## Multiclass Logistic Regression - Grayscale



# Model Experiment Process

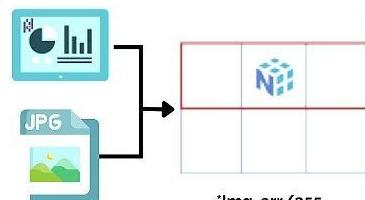




1. Choose only images of size  $150 \times 150 \times 3$



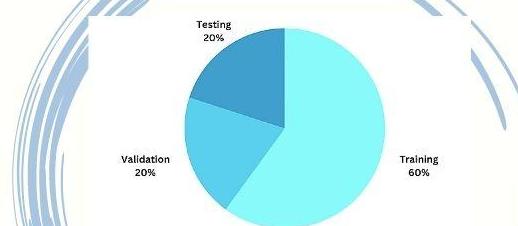
2. Generate an equal number of images from each class



3. Convert list of labels and images to numpy arrays



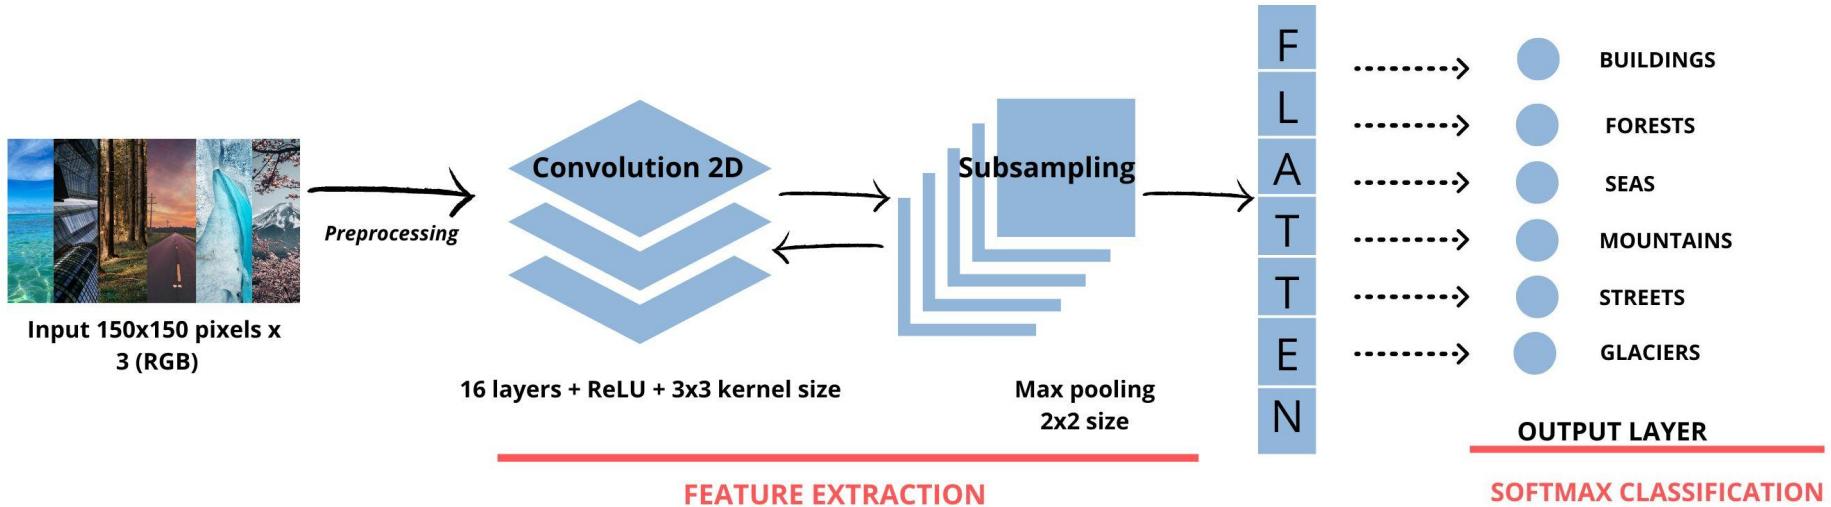
5. Shuffle the images and labels together.



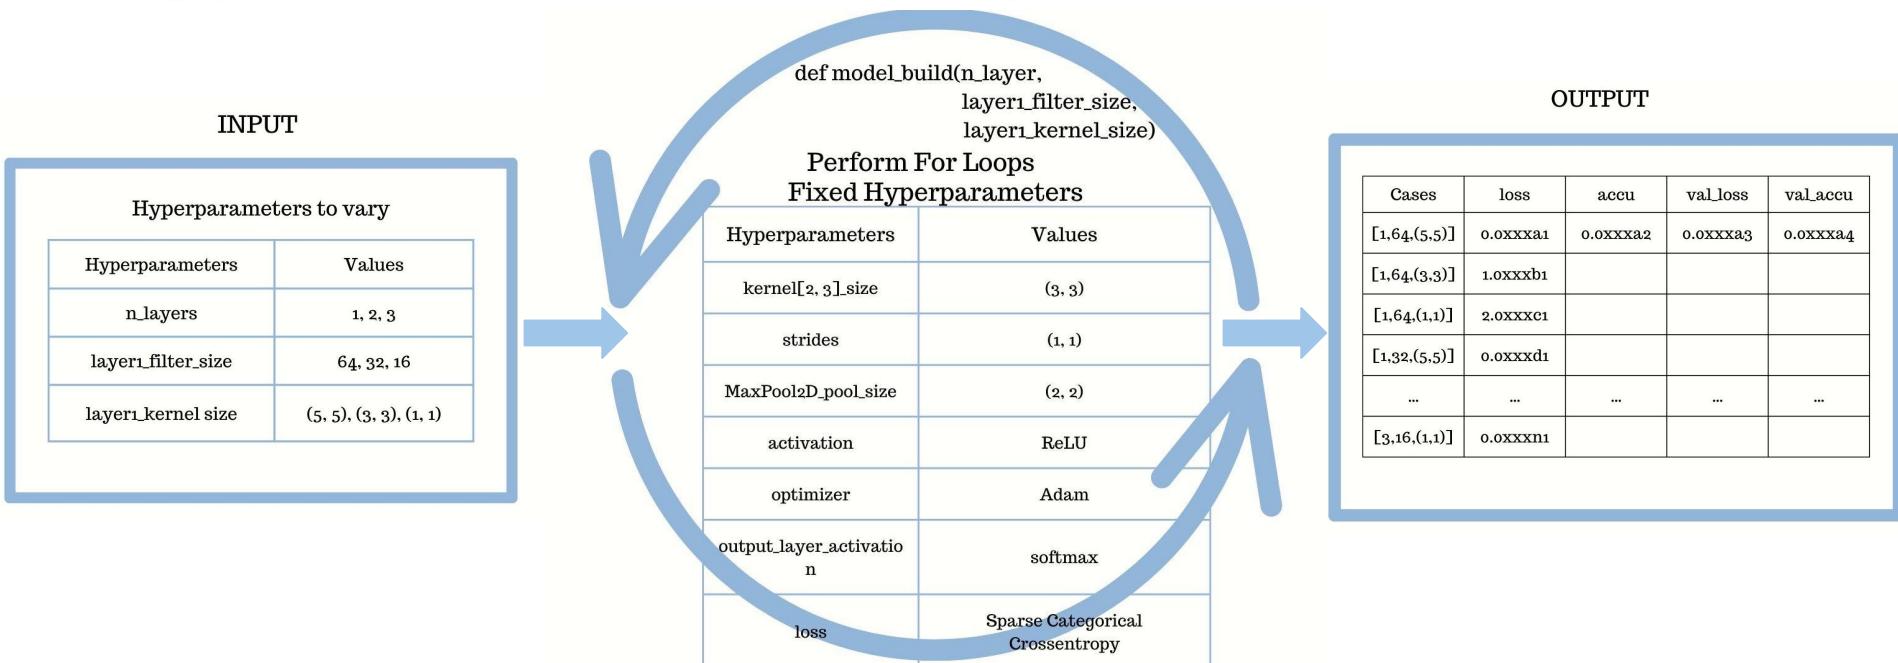
4. Split the images to comprise 60% training, 20% validation and 20% testing

## Data Preparation and Transformation

# Convolutional Neural Network 2D



# Hyperparameter Tuning



# Model Retrain: Apply CNN\_2D Optimum Case

	train_loss	train_accu	val_loss	val_accu	filter1_size	kernel1_size	num_conv_layers	build_minutes
0	0.0031	0.9999	0.0361	0.9905	64.0	(5, 5)	1	8.879024
1	0.0058	0.9996	0.0347	0.9911	32.0	(5, 5)	1	4.789623
2	0.0077	0.9997	0.0368	0.9895	16.0	(5, 5)	1	3.473544
3	0.0033	0.9998	0.0263	0.9930	64.0	(3, 3)	1	6.649201
4	0.0040	0.9998	0.0288	0.9924	32.0	(3, 3)	1	4.994755
5	0.0081	0.9998	0.0299	0.9933	16.0	(3, 3)	1	2.412871
6	0.0824	0.9823	0.1419	0.9705	64.0	(1, 1)	1	6.837631
7	0.1238	0.9722	0.1747	0.9559	32.0	(1, 1)	1	3.754270
8	0.1947	0.9533	0.2762	0.9194	16.0	(1, 1)	1	2.064861
9	0.0596	0.9861	0.0870	0.9832	64.0	(5, 5)	2	13.014805

# CNN\_2D Model Retrain at optimum HPs over more epochs

HYPERPARAMETERS	VALUES
Epochs	<b>10</b> [5, 10, 15]
No. of CNN Layers	<b>1</b> [1, 2, 3]
Filter Size	<b>64</b> [64, 32, 16]
Kernel Size	<b>(3, 3)</b> [(5, 5), (3, 3), (1, 1)]
MaxPool2D Size	(2, 2)
Conv2D activation	<i>ReLU</i>
Output Activation	<i>Softmax</i>
Optimizer	<i>Adam</i>
Loss	SCC

	loss	accuracy	val_loss	val_accuracy
<b>0</b>	1.242730	0.703500	0.318093	0.920368
<b>1</b>	0.130304	0.969546	0.085889	0.981599
<b>2</b>	0.026977	0.996511	0.052284	0.991434
<b>3</b>	0.008711	0.999471	0.027462	0.992703
<b>4</b>	0.003313	0.999789	0.026319	0.993020
<b>5</b>	0.001790	1.000000	0.036934	0.991434
<b>6</b>	0.001181	1.000000	0.026425	0.993338
<b>7</b>	0.000731	1.000000	0.026440	0.993655
<b>8</b>	0.000528	1.000000	0.026746	0.993338
<b>9</b>	0.000399	1.000000	0.025375	0.995241

# Analyze Results: CNN\_2D Classification Report

There are no indication of bias in the classes, 0 - 5. They are tightly grouped together between > 0.98 and < 1 in all metrics.

## Classes

- 0 - Buildings
- 1 - Forests
- 2 - Glacier
- 3 - Mountains
- 4 - Sea
- 5 - Street

	precision	recall	f1-score	support
0	0.9978	0.9803	0.9890	458.0000
1	1.0000	1.0000	1.0000	508.0000
2	0.9944	0.9832	0.9888	537.0000
3	0.9852	0.9962	0.9907	533.0000
4	0.9894	1.0000	0.9947	466.0000
5	0.9908	0.9954	0.9931	651.0000
<b>accuracy</b>	0.9927	0.9927	0.9927	0.9927
<b>macro avg</b>	0.9929	0.9925	0.9927	3153.0000
<b>weighted avg</b>	0.9927	0.9927	0.9927	3153.0000

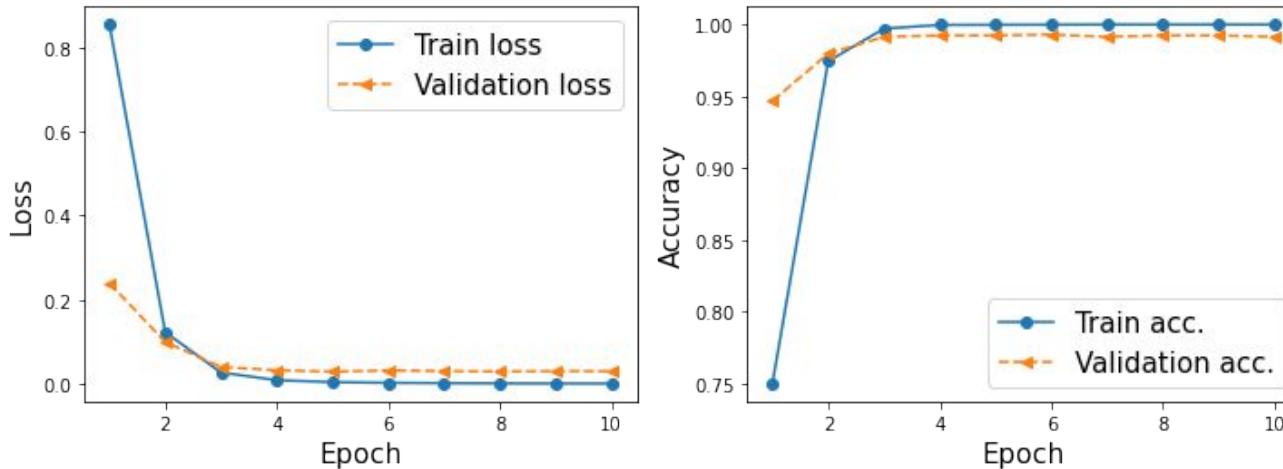
**Precision** - Percentage of correct positive predictions relative to total positive predictions.

**Recall** - Percentage of correct positive predictions relative to total actual positives.

**F1-score** - A weighted harmonic mean of precision and recall:  
$$2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$$

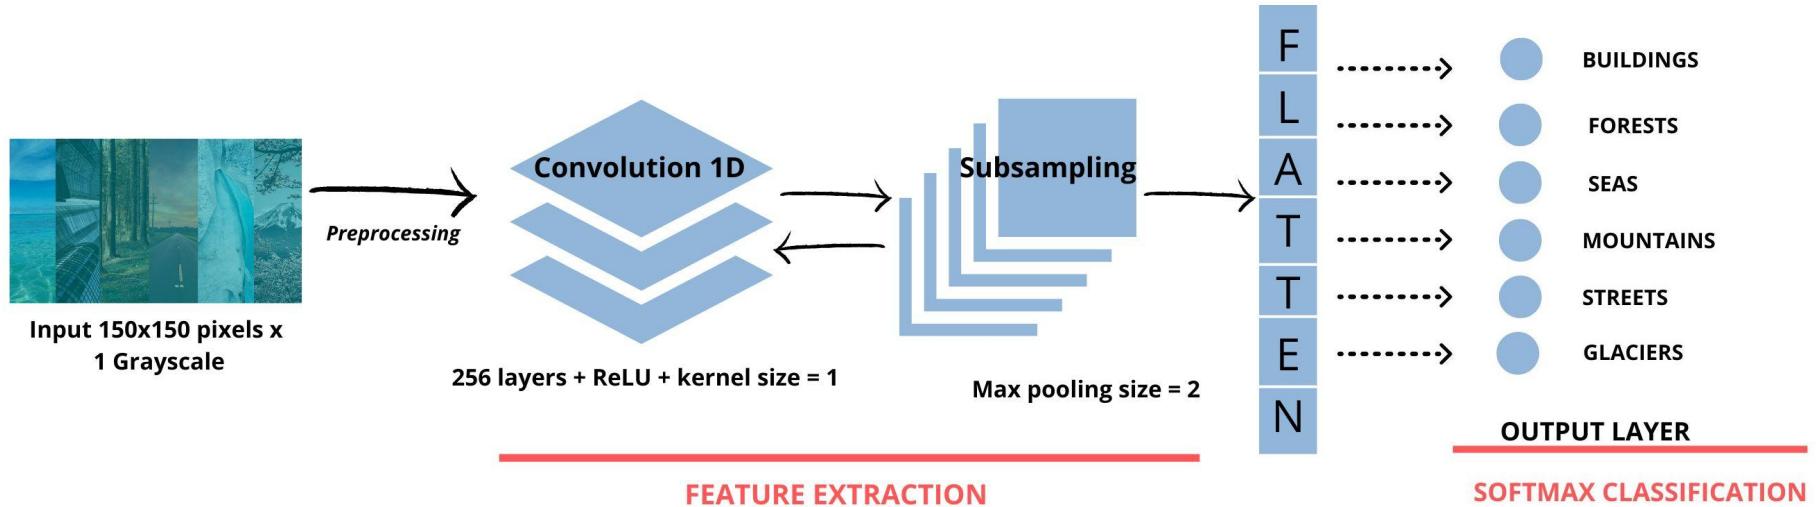
**Support** - These values tell us how many images belong to each class in the test dataset.

# Analyze Results: CNN\_2D Loss and Accuracy Plots



There are no indications of over-fitting as can be seen above where the plots converged almost from the start.

# Convolutional Neural Network 1D



# CNN\_1D: Hypertune and Apply Optimum Case

	train_loss	train_accu	val_loss	val_accu	filter1_size	kernel1_size	build_minutes
0	1.0878	0.6101	1.0513	0.6164	512.0	7	1.086736
1	1.0349	0.6254	0.9917	0.6304	256.0	7	0.785027
2	1.1566	0.5726	1.1152	0.6028	128.0	7	0.524590
3	1.0901	0.5986	1.0910	0.5945	64.0	7	0.366595
12	0.1582	0.9662	0.1822	0.9578	512.0	1	0.568469
13	0.1595	0.9712	0.1890	0.9626	256.0	1	0.362797
14	0.2300	0.9508	0.2843	0.9334	128.0	1	0.233020
15	0.3917	0.9005	0.4112	0.8877	64.0	1	0.214232

# CNN\_1D Model Retrain at optimum HPs over more epochs

HYPERPARAMETERS	VALUES
Epochs	<b>15</b> [5, 10, 15, 20]
No. of CNN Layers	<b>1</b> [1, 2, 3]
Filter Size	<b>256</b> [512, 256, 128, 64]
Kernel Size	<b>1</b> [7, 5, 3, 1]
MaxPool1D Size	2
Conv2D activation	<i>ReLU</i>
Output Activation	<i>softmax</i>
Optimizer	<i>Adam</i>
Loss	SCC

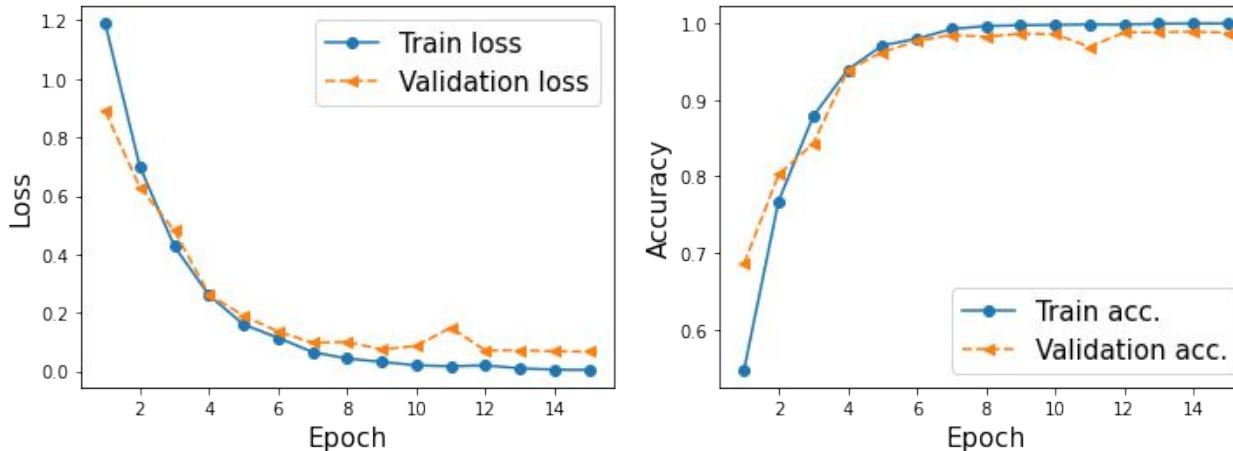
	loss	accuracy	val_loss	val_accuracy
<b>0</b>	1.189192	0.546791	0.888876	0.685596
<b>1</b>	0.699524	0.767685	0.628770	0.802348
<b>2</b>	0.426931	0.878503	0.480725	0.842957
<b>3</b>	0.259817	0.939516	0.262416	0.938769
<b>4</b>	0.159511	0.971238	0.188955	0.962563
<b>5</b>	0.114062	0.980015	0.135424	0.976840
<b>6</b>	0.065054	0.992704	0.096758	0.985089
<b>7</b>	0.042953	0.996616	0.099908	0.982551
<b>8</b>	0.031719	0.997674	0.074131	0.986675
<b>9</b>	0.019681	0.998202	0.086057	0.986041
<b>10</b>	0.016538	0.998731	0.149400	0.969226
<b>11</b>	0.019724	0.998625	0.070287	0.988261
<b>12</b>	0.009439	0.999683	0.072023	0.988579
<b>13</b>	0.005027	1.000000	0.067441	0.989213
<b>14</b>	0.003959	1.000000	0.066436	0.987944

# Analyze Results: CNN\_1D Classification Report

The results here are only slightly poorer than in CNN\_2D.  
However, the grouping is still tight and with no class showing any obvious *bias* over another.

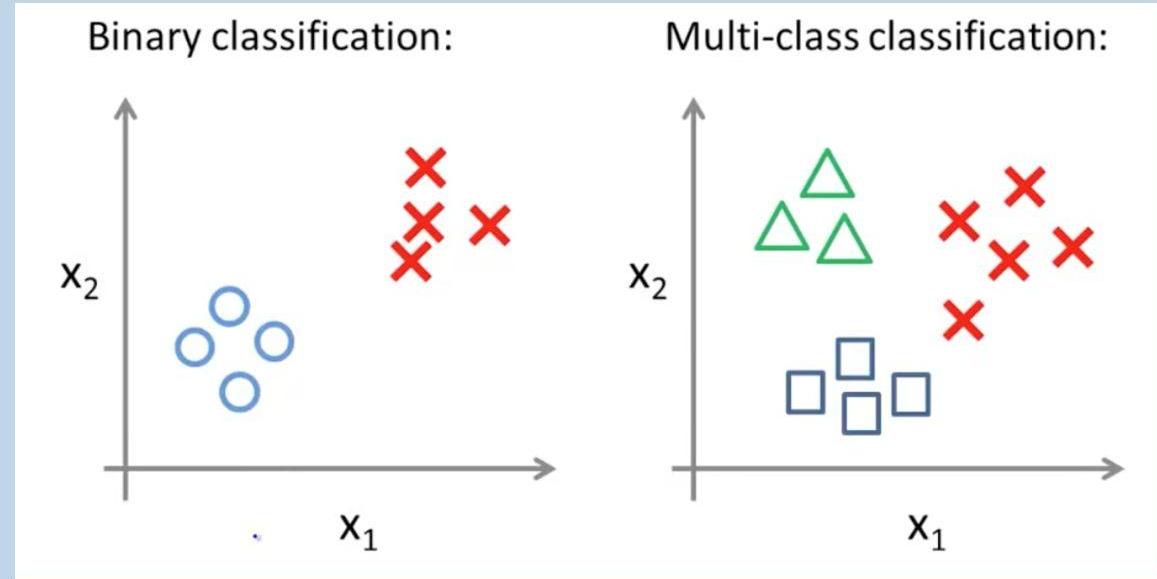
	precision	recall	f1-score	support
<b>0</b>	0.9890	0.9847	0.9869	458.0000
<b>1</b>	0.9941	0.9882	0.9911	508.0000
<b>2</b>	0.9814	0.9832	0.9823	537.0000
<b>3</b>	0.9906	0.9887	0.9897	533.0000
<b>4</b>	0.9786	0.9807	0.9796	466.0000
<b>5</b>	0.9847	0.9908	0.9877	651.0000
<b>accuracy</b>	0.9864	0.9864	0.9864	0.9864
<b>macro avg</b>	0.9864	0.9861	0.9862	3153.0000
<b>weighted avg</b>	0.9864	0.9864	0.9864	3153.0000

# Analyze Results: CNN\_1D Loss and Accuracy Plot



The plots here also show a converging pattern. But unlike in CNN\_2D, the validation line show a bit of a bump at the 3rd and 11th epoch.

# Multi-class Logistic Regression- Grayscale



# MCLR Model Retrain at optimum HPs over more epochs

HYPERPARAMETERS	VALUES
Epochs	<b>50</b> [10 - 60, step: 10]
Learning Rate	<b>0.001</b> [0.01, 0.001]
Batch Size	<b>32</b> [16, 32, 64]
Optimizer	<b>Adam</b> [SGD, Adam]
Output Activation	<b>softmax</b>
Loss	<i>Sparse Categorical Crossentropy (SCC)</i>

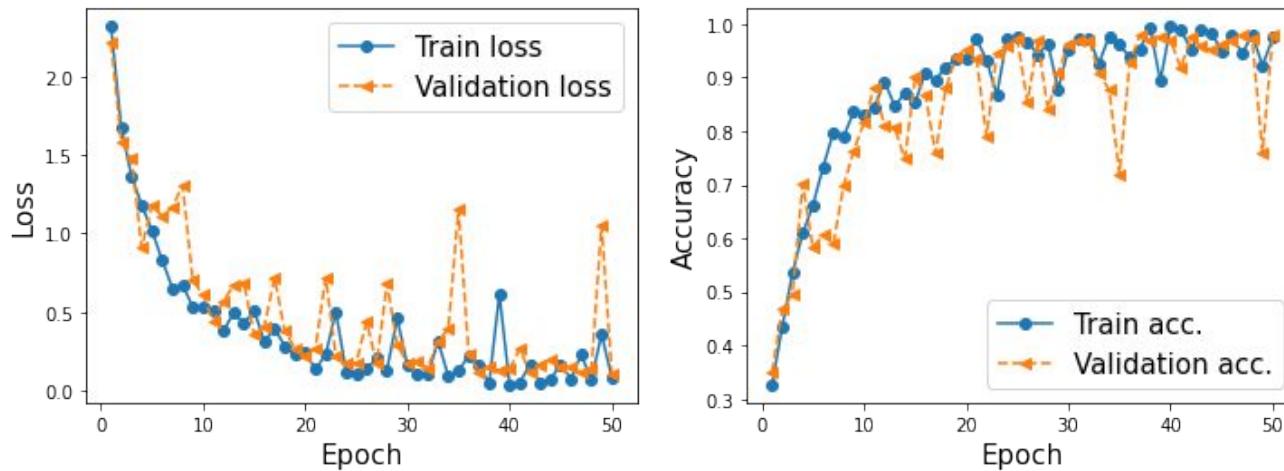
	loss	accuracy	val_loss	val_accuracy
<b>0</b>	2.321057	0.326319	2.214630	0.348985
<b>1</b>	1.672324	0.434387	1.581540	0.468274
<b>2</b>	1.367396	0.534842	1.484961	0.494607
<b>3</b>	1.180820	0.609073	0.907844	0.701777
<b>4</b>	1.017916	0.661521	1.179285	0.582487
<b>45</b>	0.072881	0.979380	0.146761	0.968274
<b>46</b>	0.227095	0.944697	0.113168	0.977792
<b>47</b>	0.064674	0.980226	0.132256	0.971129
<b>48</b>	0.355281	0.921857	1.055452	0.758883
<b>49</b>	0.083987	0.975256	0.108018	0.980013

# Analyze Results: MCLR Classification Report

Although MCLR is the worst of the three in terms of loss and accuracy, the classification report still show no indication of bias and it still achieved a fairly decent accuracy score of 0.98.

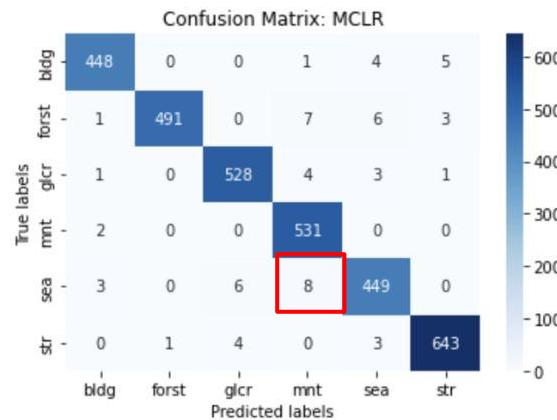
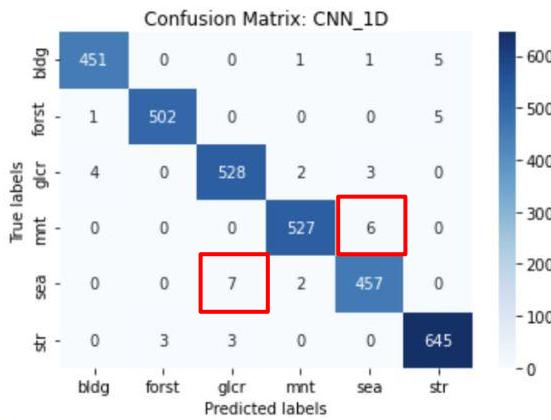
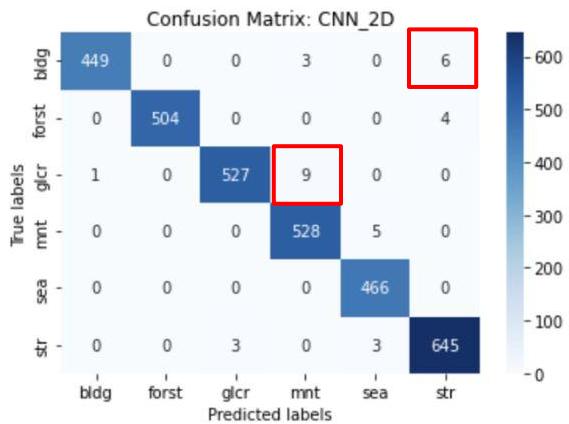
	precision	recall	f1-score	support
<b>0</b>	0.9846	0.9782	0.9814	458.00
<b>1</b>	0.9980	0.9665	0.9820	508.00
<b>2</b>	0.9814	0.9832	0.9823	537.00
<b>3</b>	0.9637	0.9962	0.9797	533.00
<b>4</b>	0.9656	0.9635	0.9646	466.00
<b>5</b>	0.9862	0.9877	0.9870	651.00
<b>accuracy</b>	0.9800	0.9800	0.9800	0.98
<b>macro avg</b>	0.9799	0.9792	0.9795	3153.00
<b>weighted avg</b>	0.9802	0.9800	0.9800	3153.00

# Analyze Results: Multi-class Logistic Regression (MCLR) Loss and Accuracy Plot



The path taken to achieve a  $> 0.98$  score is erratic as shown above.

# Comparing Confusion Matrices



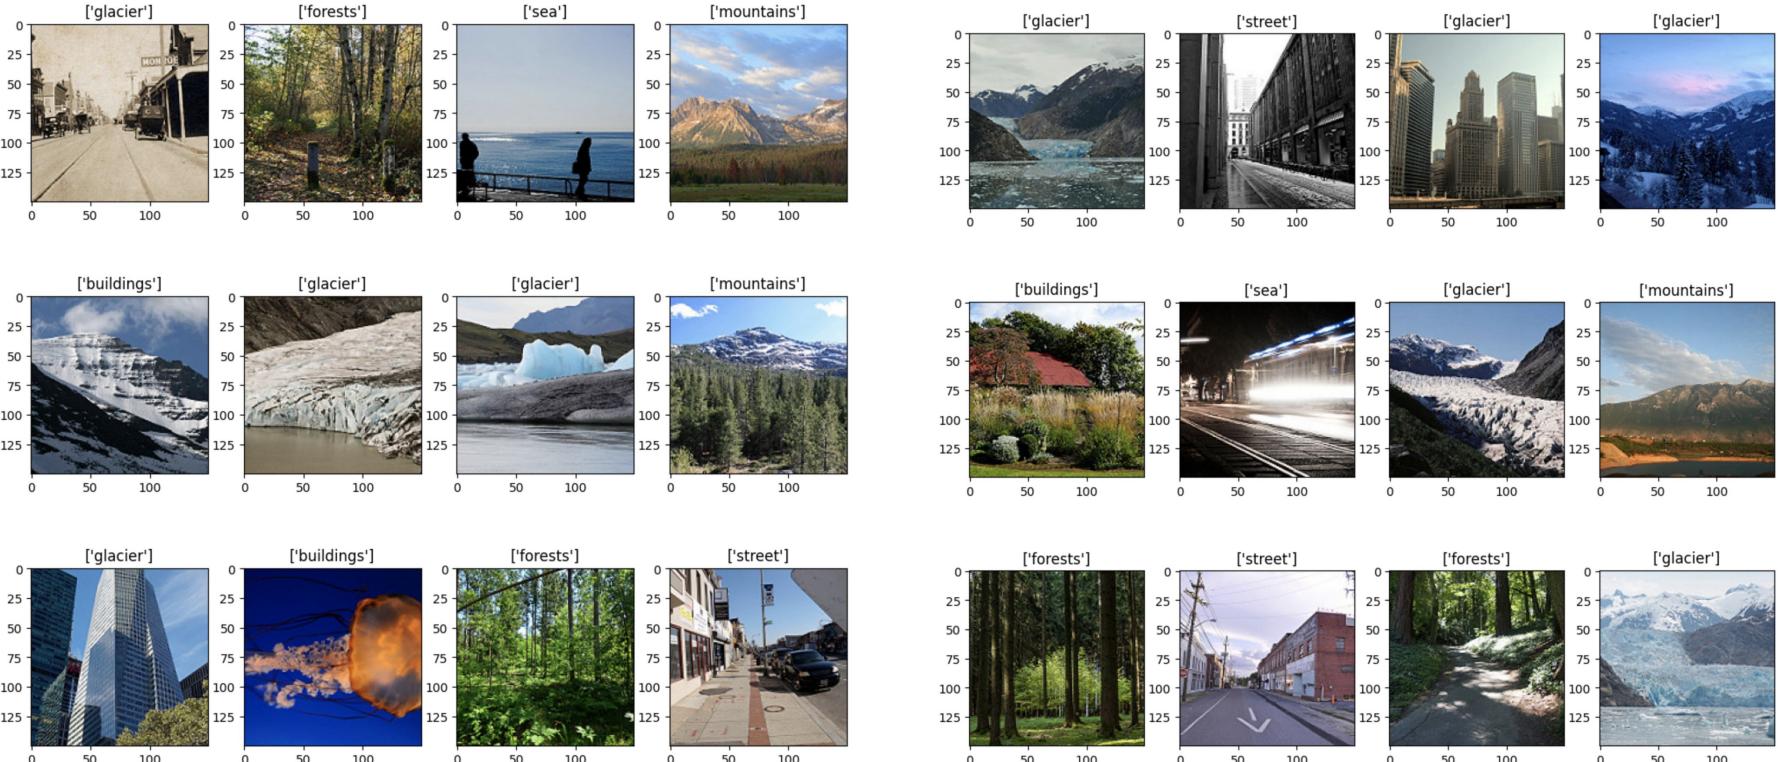
The confusion matrices vary in results. CNN\_2D errors were mostly mountains vs. glaciers and street vs. buildings; CNN\_1D: sea/glaciers and sea/mountain; MCLR: sea/mountain and forest/mountain

# Conclusion: Model vs. Test Data

For our use case, model training will occur offline so training times don't matter. Therefore, due to its slightly higher classification accuracy, we would recommend a 2D CNN for classifying images in a production environment.

<i>model.evaluate()</i>	Loss	Accuracy
CNN_2D	0.030794	0.992705
CNN_1D	0.066340	0.986362
MCLR	0.150250	0.980019

# Conclusion: Final Test



Thank you.